

HEPDataset: A configurable framework for constructing high dimensional machine learning datasets from Delphes ROOT files

Yusheng Wu¹ and A. M. M Elsayed¹

¹ Department of Modern Physics and State Key Laboratory of Particle Detection and Electronics, University of Science and Technology of China, Hefei, China  Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

The search for physics beyond the Standard Model at the Large Hadron Collider increasingly uses supervised and unsupervised machine learning to distinguish rare signal processes from substantially larger Standard Model backgrounds. Phenomenological studies commonly begin by generating Monte Carlo samples and then simulating the detector response with Delphes fast detector simulation (Favereau et al., 2014). The expected input to the current HEPDataset workflow is therefore a ROOT file containing Delphes events. Such analyses require event level representations that preserve the information contained in reconstructed detector objects. The construction of these representations from Delphes output is often performed by an analysis specific event loop. This task requires repeated access to nested ROOT objects, object quality selection, kinematic feature construction, analysis channel classification, sample normalization, and output management. These responsibilities are closely coupled in many analysis scripts. The resulting workflow is time consuming to develop and difficult to audit for consistency.

HEPDataset is an open source software package that provides a configurable framework for reading Delphes ROOT files and producing machine learning ready event datasets. The package uses a branch configuration card named `branches_config.yml` to define the reconstructed objects, object representations, multiplicities, kinematic observables, event level quantities, and multi object combinations that should be extracted. A single configuration can request a high dimensional feature space containing up to approximately 12000 branches, subject to the available content of the input samples and the resources of the processing environment. The configuration therefore allows an analysis to preserve a broad description of each event before applying feature reduction or model specific selection.

The package performs object selection and analysis channel classification during the same event traversal that computes the requested features. It can therefore classify an event and populate all relevant output representations during one pass through the input file. This design avoids repeated reading of the same ROOT data when different channels or derived observables are required. The analysis channel description is configurable through `analysis_channels.py`. The object selection policy is configurable through `object_selection.py` and is intended to become more declarative through dedicated YAML cards in future releases. The current loop implementations use PyROOT to access the input ROOT files. Alternative loop implementations based on Uproot are planned so that the same dataset construction model can be used through a more Python oriented interface. These alternatives are complementary because Uproot is intended as another input backend rather than as a competitor to HEPDataset.

The workflow accepts collections of signal and background samples with their associated cross sections. It derives event weights using the generated event count, the process cross section, and the target integrated luminosity. It writes channel specific ROOT datasets and

can merge outputs from different samples and processing chunks. Event processing is divided into independent ranges of the input event sequence. Each range is written to an isolated temporary output. The temporary outputs are then merged into the final dataset. This design permits parallel execution without concurrent writes to a common ROOT file.

HEPDataset is intended primarily to make the rapid generation of high dimensional machine learning datasets possible. The package can support conventional analysis workflows that use tabular classifiers and it can also support the future derivation of expected discovery significance for BSM signal hypotheses. Its principal goal is broader. It is intended to provide large and information rich event representations that can be supplied to advanced machine learning frameworks, including architectures based on attention and transformer mechanisms. Such models may benefit from access to thousands of correlated and complementary observables rather than to a small set of variables selected in advance by manual analysis practice.

Statement of need

Delphes provides a widely used fast simulation of a generic collider detector (Favreau et al., 2014). Phenomenological studies frequently use it after Monte Carlo event generation because it provides a computationally efficient approximation of the detector response. The current expected input to HEPDataset is a ROOT file produced by Delphes. Other detector simulation systems, including Geant4 and ATLASFast, can be supported in future releases provided that their event information is available in a compatible ROOT based representation. ROOT provides efficient storage and access to these data (Brun & Rademakers, 1997). The structure is appropriate for high energy physics analysis, but it is not directly equivalent to a machine learning table. The information for one event is distributed across object collections with event dependent multiplicity. A machine learning application generally requires a stable feature representation with explicit treatment of missing objects, object ordering, derived quantities, and event weights.

The direct construction of such a representation is time consuming because a researcher must repeatedly translate a physics definition into data access operations. Each new observable can require a new branch declaration, a new object traversal, a new convention for missing values, and a new validation procedure. The difficulty is not limited to the arithmetic of a kinematic variable. It also concerns the consistent treatment of object selection, combinatorial assignments, flavour categories, charge categories, signal regions, and sample normalization. When these responsibilities are implemented separately for multiple analyses, small inconsistencies can propagate into the training data and make comparisons between signal and background samples less reliable.

The direct approach is also vulnerable to avoidable errors. An analysis can accidentally read different object collections for different regions. It can calculate a variable only for some channels. It can apply incompatible selection criteria to signal and background samples. It can normalize samples using an incorrect event count or unit convention. It can read the same ROOT file several times because the output requirements for different regions were developed independently. These issues are particularly consequential for rare signal searches because a small selection or normalization discrepancy can alter the apparent separation between signal and background.

HEPDataset addresses this need by separating the definition of the desired physics representation from the repetitive mechanics of event processing. The `branches_config.yml` card provides a declarative description of the feature space. It can define reconstructed objects, their representations, the number of objects to retain, the kinematic quantities to derive, and the combinations of objects that should be considered. The card can therefore describe a feature space with up to approximately 12000 candidate features while retaining a single source of truth for the branch specification. The resulting representation can contain single object kinematics, multi object invariant quantities, angular correlations, missing transverse

93 momentum observables, global activity variables, and event shape observables.

94 The feature space can include the kinematic content associated with reconstructed four vectors,
95 including transverse momentum, pseudorapidity, azimuthal angle, invariant mass, energy,
96 and the Cartesian momentum components. It can also include tracking and reconstruction
97 parameters for leptons, jets, and fat jets whenever those quantities are available in the input
98 ROOT branches. For pairwise and higher order combinations, the package can derive angular
99 and transverse observables such as ΔR , $\Delta\eta$, $\Delta\phi$, invariant mass, and M_{T2} , with M_{T2} defined
100 for the relevant two object and missing transverse momentum configuration. It can construct
101 additional kinematic quantities from vector sums and scalar sums of object four vectors.
102 Event level information can include missing transverse momentum, hadronic activity, lepton
103 activity, total activity, effective mass, missing momentum significance, and the angular and
104 kinematic properties of the missing momentum. Event shape observables such as centrality,
105 circularity, aplanarity, and sphericity can be evaluated for the full event or for selected groups
106 of reconstructed particles. These categories illustrate how one event can provide many single
107 object features, an even larger set of pairwise and higher order features, and a broad collection
108 of global event variables.

109 The current Python event loops use the PyROOT framework to read the ROOT files and
110 access the Delphes event content. The package already provides alternative loop choices with
111 different levels of explicitness and configurability. For example, a fixed loop can access a lepton
112 transverse momentum through the direct attribute `lep.PT`, while the adaptive loop can obtain
113 the same property through a configurable attribute lookup. This distinction allows the adaptive
114 approach to expose a larger feature space through a compact configuration card. Future
115 releases will add loop implementations based on Uproot, including counterparts analogous to
116 `basic3_uproot.py` and `adaptive_uproot.py`. Uproot is not considered a competing package.
117 It is a Python oriented ROOT input backend that can make the data access layer more
118 independent of the ROOT runtime. The repository also contains C++ loop infrastructure
119 intended for future serial and parallel ROOT processing. Each loop implementation is designed
120 to accept a ROOT file path and return either channel keyed trees, a channel keyed dictionary
121 of trees, or channel keyed paths to written ROOT files.

122 Future versions will also incorporate advanced learned taggers and related scoring models into
123 the feature construction stage. Models in the ParticleNet family and comparable architectures
124 can provide scores for b tagged, tau tagged, and c tagged jets. Dedicated models can also
125 provide scores for QCD fat jets, W fat jets, Z fat jets, Higgs fat jets, and top fat jets. These
126 scores can be recorded as additional event features alongside the conventional kinematic and
127 reconstruction quantities. The intended objective is to read each ROOT event once and to
128 produce a broad collection of complementary observables and learned object scores that are
129 immediately available for downstream analysis.

130 The package also performs analysis channel selection during the event traversal. The same
131 selected objects that provide the feature values determine the event category. This permits one
132 pass through a ROOT file to produce the information required for multiple analysis channels.
133 The approach is intended to be straightforward for researchers because changes to the feature
134 space can be expressed through the configuration card rather than through a complete rewrite
135 of the data extraction workflow. It is intended to be efficient because event access, selection,
136 feature construction, and channel assignment are combined within one processing pass. It is
137 intended to be extensible because the channel and object selection descriptions can evolve
138 independently of the sample registry and the processing strategy.

139 The package is designed for analysis groups that need to construct datasets before a final
140 machine learning architecture has been selected. A broad candidate feature space allows later
141 studies to compare classical feature engineering with models that can learn representations from
142 high dimensional inputs. The package therefore addresses a practical gap between detector
143 simulation and advanced machine learning rather than replacing event generators, detector
144 simulation software, recasting frameworks, or statistical inference tools.

145 The current workflow is summarized in Figure 1.

flowchart TD

```

A[Delphes ROOT samples] --> B[Sample registry]
C[Branch configuration card] --> D[Feature space definition]
E[Object selection configuration] --> F[Selected reconstructed objects]
G[Analysis channel configuration] --> H[Channel assignment]
B --> I[Single event traversal]
D --> I
F --> I
H --> I
A --> I
I --> J[Up to approximately 12000 event features]
I --> K[Cross section and luminosity weights]
I --> L[Channel specific datasets]
J --> L
K --> L
L --> M[Machine learning and physics analysis]

```

146 State of the field

147 HEPDataset addresses the event representation stage of the analysis lifecycle. Its input is
 148 a ROOT file containing simulated or reconstructed events. The package can use PyROOT,
 149 Uproot, and, in future releases, C++ based event loops to traverse the input events and extract
 150 the requested information. Its target use case is the construction of event level datasets for
 151 an analysis that may use a multivariate classifier and whose feature space may evolve during
 152 the research process. The central output is not only a yield in a predefined signal region. It is
 153 a reproducible and weighted representation of many individual events with a large collection
 154 of physics motivated features. The event loop can be divided into independent ranges and
 155 executed in parallel before the resulting channel organized ROOT files are merged. This output
 156 can be used to train a BDT, a neural network, or a future architecture that learns from high
 157 dimensional event information.

158 The distinction is important for BSM searches. Recasting generally starts with an analysis whose
 159 object definitions and selection regions are already specified. Machine learning based analysis
 160 development can start earlier. The researcher may wish to retain a wide set of observables and
 161 determine through statistical learning which combinations contain useful information about a
 162 signal hypothesis. The researcher may also wish to compare several signal regions, mass points,
 163 background mixtures, and feature definitions without rebuilding the entire data production
 164 chain.

165 HEPDataset is therefore complementary to existing HEP software. It does not attempt to
 166 reproduce the detector simulation provided by Delphes. It does not attempt to replace the
 167 analysis validation and recasting capabilities of MadAnalysis 5, CheckMATE, or Rivet. It
 168 provides a configurable data preparation layer that can supply downstream machine learning
 169 frameworks and can later be connected to statistical tools such as pyhf (Heinrich et al., 2021).

170 The package is also complementary to generic machine learning libraries. XGBoost provides
 171 scalable gradient boosted decision trees (Chen & Guestrin, 2016). HEPDataset does not
 172 provide a replacement for XGBoost. It provides a physics aware method for constructing the
 173 input representation that a classifier requires. Future integrations with attention based and
 174 transformer based architectures can extend this role to models that operate on high dimensional
 175 or structured event representations.

176 The following table summarizes the intended division of responsibility.

Software layer	Primary responsibility	Role of HEPDataset
Event generation	Production of hard scattering events and decays	External input
Detector simulation	Approximate detector response and reconstruction	External input from Delphes
Event dataset construction	Selection, feature extraction, weighting, and organization	Primary role
Machine learning	Classification, representation learning, and ranking	Downstream or planned integration
Statistical inference	Significance estimation and hypothesis testing	Downstream or planned integration

177 Software design

178 HEPDataset is designed around a separation between the physics description of the event
 179 representation and the operational process that constructs the dataset. The user specifies the
 180 desired objects, observables, channels, regions, and samples. The package then applies those
 181 definitions consistently while handling event traversal, weighting, parallel execution, temporary
 182 output, and merging. The loop interface is deliberately centered on the input ROOT file
 183 and the channel organized result. A loop can return in memory trees, a dictionary whose
 184 keys identify analysis channels, or paths to ROOT files that have already been written. This
 185 permits PyROOT based, future Uproot based, and future C++ based loop implementations to
 186 participate in the same higher level dataset workflow.

187 The branch configuration card is the main interface for defining the candidate feature space.
 188 It allows the user to express the physical content of the dataset without repeatedly modifying
 189 the event reading workflow. The card can request object level quantities and multi object
 190 quantities. It can also define the combinations of objects used for derived variables. This
 191 is important because many discriminative properties of collider events are relational. The
 192 invariant mass of two objects, their angular separation, and their transverse momentum balance
 193 can contain information that is not represented by the individual object values alone.

194 The package reads the selected event information and assigns an analysis channel during the
 195 same traversal. The channel representation can distinguish lepton multiplicity, lepton flavour,
 196 charge structure, and jet content. The configuration in `analysis_channels.py` allows the
 197 channel taxonomy to follow the needs of a particular physics analysis. The selection policy in
 198 `object_selection.py` provides the corresponding object quality requirements. Future releases
 199 will extend this configurability through dedicated YAML cards so that object definitions and
 200 selection criteria can be recorded as analysis data rather than only as program logic.

201 The single traversal design is motivated by both computational and scientific considerations.
 202 Repeated traversal of a large ROOT file increases input output cost and creates additional
 203 opportunities for inconsistent selection. A single traversal ensures that the features, event
 204 category, and event weight are derived from the same reconstructed event state. It also permits
 205 all requested channels to be produced from the same input pass. The approach is particularly
 206 valuable when the feature space contains thousands of candidate variables.

207 The processing strategy uses event range partitioning. A ROOT input file is divided into
 208 independent intervals of entries. Each worker receives one interval and produces isolated
 209 temporary outputs. The workers do not write to a common ROOT file. The resulting files are
 210 merged after the workers complete. The isolation of writes avoids concurrent access conflicts
 211 and makes the worker computation close to embarrassingly parallel. It also provides a clear
 212 failure boundary because a failed interval can be identified without corrupting a shared output
 213 file.

214 The merging strategy follows the structure of the physics output. Temporary files are first
215 merged within a sample and channel or region. Sample level outputs can then be merged
216 into final datasets for the corresponding channel or region. This hierarchy preserves the
217 provenance of intermediate products while allowing the user to retain per sample outputs or
218 create combined training datasets. Temporary files are removed after a successful merge so
219 that the storage requirement does not grow with the number of processing intervals.

220 The complete current processing sequence is shown in Figure 2.

sequenceDiagram

```

participant U as Researcher
participant C as Configuration cards
participant P as Dataset pipeline
participant R as Delphes ROOT file
participant W as Parallel workers
participant H as ROOT merger
participant O as ML dataset

U->>C: Define samples objects features and channels
U->>P: Start dataset construction
P->>C: Validate and load configuration
P->>R: Open one input sample
P->>W: Assign independent event ranges
W->>R: Read assigned events
W->>W: Select objects compute features and assign channels
W->>W: Write isolated temporary ROOT outputs
W-->>P: Return temporary output locations
P->>H: Merge temporary outputs
H-->>P: Return sample level channel outputs
P->>H: Merge selected samples when requested
H-->>O: Write final channel datasets

```

221 The design favors a transparent and configurable data production process over a fully automatic
222 interpretation of physics intent. The package can calculate requested observables and apply
223 requested selections. It cannot determine whether a particular selection is theoretically optimal
224 or whether a variable is physically appropriate for a given search. That responsibility remains
225 with the analysis team and requires validation against the detector card, the signal model, and
226 independent physics calculations.

227 The package is intended to support two complementary modes of use. In the first mode, a
228 researcher generates a broad dataset and supplies it to an external machine learning framework.
229 In the second mode, a future integrated workflow selects a smaller set of informative features
230 and trains a classifier within the same analysis environment. The second mode is planned
231 rather than a claim about the complete current release.

232 Future development will introduce a discriminative feature selection stage. The user will
233 be able to define a maximum number of retained features. The selection stage will seek
234 features that provide strong discrimination between signal and background while controlling
235 redundancy through an uncorrelatedness requirement. The purpose is to reduce the high
236 dimensional candidate space to a compact representation that remains physically interpretable
237 and computationally manageable.

238 The planned machine learning and significance workflow is shown in Figure 3.

flowchart TD

```

A[High dimensional weighted dataset] --> B[Quality and validity checks]
B --> C[Signal and background partitions]
C --> D[Correlation assessment]

```

```

C --> E[Discriminative ranking]
D --> F[User defined maximum feature count]
E --> F
F --> G[Selected low redundancy feature set]
G --> H[XGBoost BDT training]
H --> I[Classifier score]
I --> J[Signal region and threshold scan]
J --> K[Expected discovery significance]
K --> L[Physics interpretation]

```

239 This future stage will complement rather than replace the high dimensional dataset generator.
 240 The broad dataset remains the important scientific product because it permits the use of
 241 advanced learning systems that can discover nonlinear and higher order relationships among
 242 many observables. Feature reduction provides an optional route for interpretable BDT based
 243 studies and for applications where computational cost or systematic control favors a smaller
 244 representation.

245 Research impact statement

246 HEPDataset is intended to reduce the time required to move from Delphes detector simulation
 247 to an analysis ready machine learning dataset. Its impact is therefore measured not only by
 248 the number of features or processed events. It is measured by whether analysis groups can
 249 construct, modify, validate, and compare event representations with less duplicated engineering
 250 effort.

251 The package can support searches for BSM physics in which the signal is distributed across
 252 several correlated kinematic properties. A conventional cut based analysis may use a small
 253 number of observables selected from prior physical intuition. HEPDataset permits the analyst
 254 to retain a substantially broader candidate space before the final classifier architecture or
 255 feature selection procedure is chosen. This is relevant for signals with complex decay chains,
 256 multiple production modes, compressed spectra, or experimentally overlapping final states.

257 The package also supports reproducible comparison between analyses. The branch configuration
 258 card records the requested feature space. The sample registry records the input processes and
 259 cross sections. The channel and object selection descriptions record the event categorization.
 260 Together these inputs provide a more explicit account of the data construction process than a
 261 collection of disconnected analysis scripts.

262 The planned feature selection and XGBoost integration will enable a complete exploratory
 263 workflow in which candidate variables are ranked, a BDT is trained, and the classifier score
 264 is used to estimate expected discovery significance for BSM signals. Such a workflow can
 265 help identify promising signal models and guide the design of later analysis selections. The
 266 significance estimate will remain dependent on the statistical model, systematic uncertainties,
 267 signal and background normalization, and assumptions about the dataset. It will therefore be
 268 an analysis aid rather than an automatic claim of discovery sensitivity.

269 The broader objective is to make the rapid generation of high dimensional event datasets routine.
 270 The resulting datasets can be supplied to advanced machine learning systems that exploit
 271 thousands of input features. Architectures based on attention and transformer mechanisms
 272 may learn relationships among event observables that are difficult to encode through manual
 273 feature selection. HEPDataset provides the data preparation foundation for these studies while
 274 preserving the possibility of using established classifiers such as XGBoost.

275 The software is intended for research use in collider phenomenology and machine learning
 276 based HEP analysis. Its outputs should be validated using independent physics calculations,
 277 controlled samples, closure tests, and systematic uncertainty studies before they are used for a
 278 physics claim. The project will benefit from future contributions that improve configuration

validation, add dedicated selection cards, test serial and parallel equivalence, expand output
format support, and connect the feature generation stage to modern high dimensional learning
architectures.

AI usage disclosure

Generative artificial intelligence was used to assist with the revision of this manuscript. The
assistance was limited to restructuring prose, improving academic language, and organizing the
description of the software according to the requested Journal of Open Source Software sections.
The repository source files, configuration examples, existing manuscript, and stated software
behavior were reviewed by the authors before the revised text was prepared. Claims concerning
current functionality were constrained to behavior represented in the repository. Planned feature
selection, classifier training, and discovery significance functionality are explicitly described as
future developments rather than as capabilities of the current release. The authors remain
responsible for the scientific accuracy, completeness, and final content of this manuscript.

Brun, R., & Rademakers, F. (1997). ROOT — an object oriented data analysis framework.
Nuclear Instruments and Methods in Physics Research Section A, 389(1-2), 81–86. [https://doi.org/10.1016/S0168-9002\(97\)00048-X](https://doi.org/10.1016/S0168-9002(97)00048-X)

Chen, T., & Guestrin, C. (2016). XGBoost: A scalable tree boosting system. *Proceedings
of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data
Mining*, 785–794. <https://doi.org/10.1145/2939672.2939785>

Favereau, J. de, Delaere, C., Demin, P., Giammanco, A., Lemaître, V., Mertens, A., &
Selvaggi, M. (2014). DELPHES 3, a modular framework for fast simulation of a generic
collider experiment. *Journal of High Energy Physics*, 2014(2), 57. [https://doi.org/10.1007/JHEP02\(2014\)057](https://doi.org/10.1007/JHEP02(2014)057)

Heinrich, L., Feickert, M., Stark, G., & Cranmer, K. (2021). Pyhf: Pure-python implementation
of HistFactory statistical models. *Journal of Open Source Software*, 6(58), 2823. <https://doi.org/10.21105/joss.02823>